

COMP 433: Introduction to Deep Learning

Practical Matters

- Labs ~10 Labs, 4 to be submitted
- 2 Problem sets
- 3 Quizzes
- Final Project — to be assigned by Week 5
- Textbook: Deep Learning by Goodfellow, Courville, Bengio — freely available online
- Dive into Deep Learning, freely available online

More Practical Matters

- Lab 1 is posted
- Office hours
 - Live office hours are Monday at 4-4:45 in room ER 958
 - By appointment on MS Teams
- MS Teams
 - Join MS Teams, link on moodle
 - Some announcements will be made there

Evaluation

- Assignments: 30%
 - Programming assignments (roughly 80% of assignments)
 - Some theoretical and written questions
- Project: 30%
- Quizzes: 30%
 - 3 Quizzes
- Labs: 10%
 - Only first 4 are graded

Pre-requisites

- Linear Algebra, Multivariable Calculus, Probability and Statistics, Algorithms
- Knowledge of *some* programming language
- Key concepts reviewed when appropriate

Outline

- Machine Learning Foundations
- Introduction to Neural Networks
- Backpropagation and Automatic Differentiation Software
- Optimization for Deep Learning
- Practical Training Recipes
- Convolutional Neural Networks
- RNNs and sequence models
- Attention and Self-Attention
- Multi-Task and Transfer Learning
- Deep Generative Models
- Self-supervised Learning

What You Will Learn From This Class


- In-depth practical and (some) theoretical understanding of the building blocks of deep learning models
- Hone or introduce important technologies like Python, Numpy, Pytorch
- Overview of some cutting edge techniques

Concepts to Reinforce or Learn

- Covered in lecture:
 - Basic linear algebra concepts
 - Relevant machine learning concepts
- From labs and additional material:
 - Python - one of the most used programming languages in the world
 - Using Tensor/Matrix libraries - e.g. Numpy

Applications

Autonomous Driving



Speech Recognition



Machine Translation




Image generation with prompts

Text Prompt

an armchair in the shape of an avocado. . . .

AI Generated images




<https://openai.com/research/dall-e>

Deep Learning

- Subset of machine learning
- The main object of study is the Neural Network
- Associated with modular but powerful/expressive framework for creating predictive models

Neural Networks

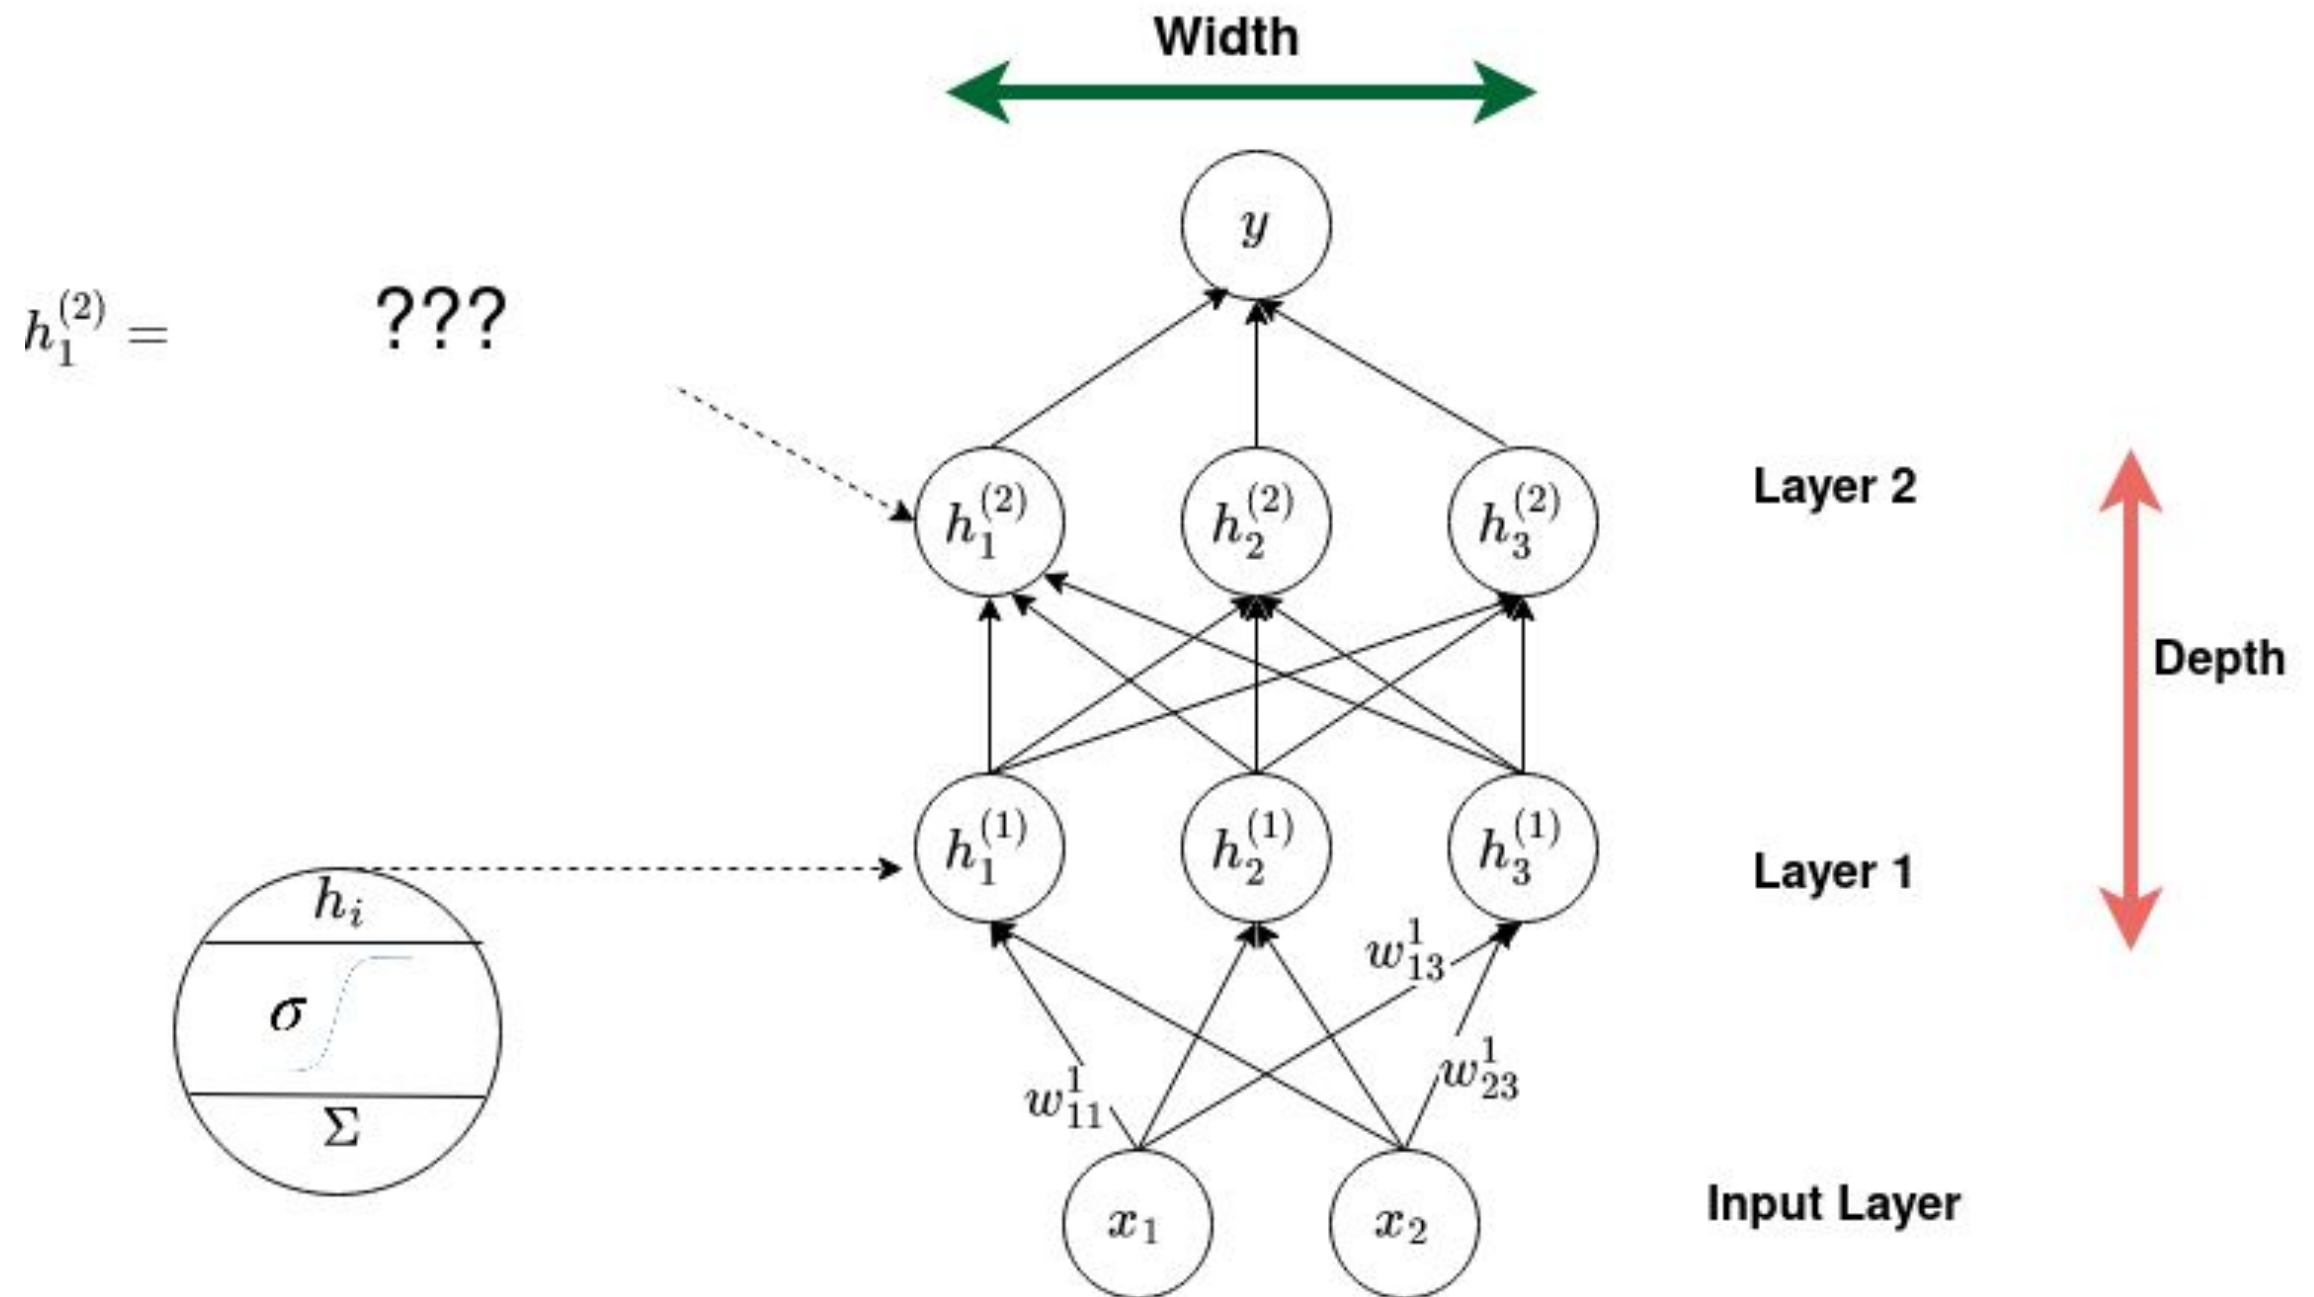
- Simple computation blocks



$$h(x_1, x_2) = \sigma(w_1x_1 + w_2x_2 + b)$$

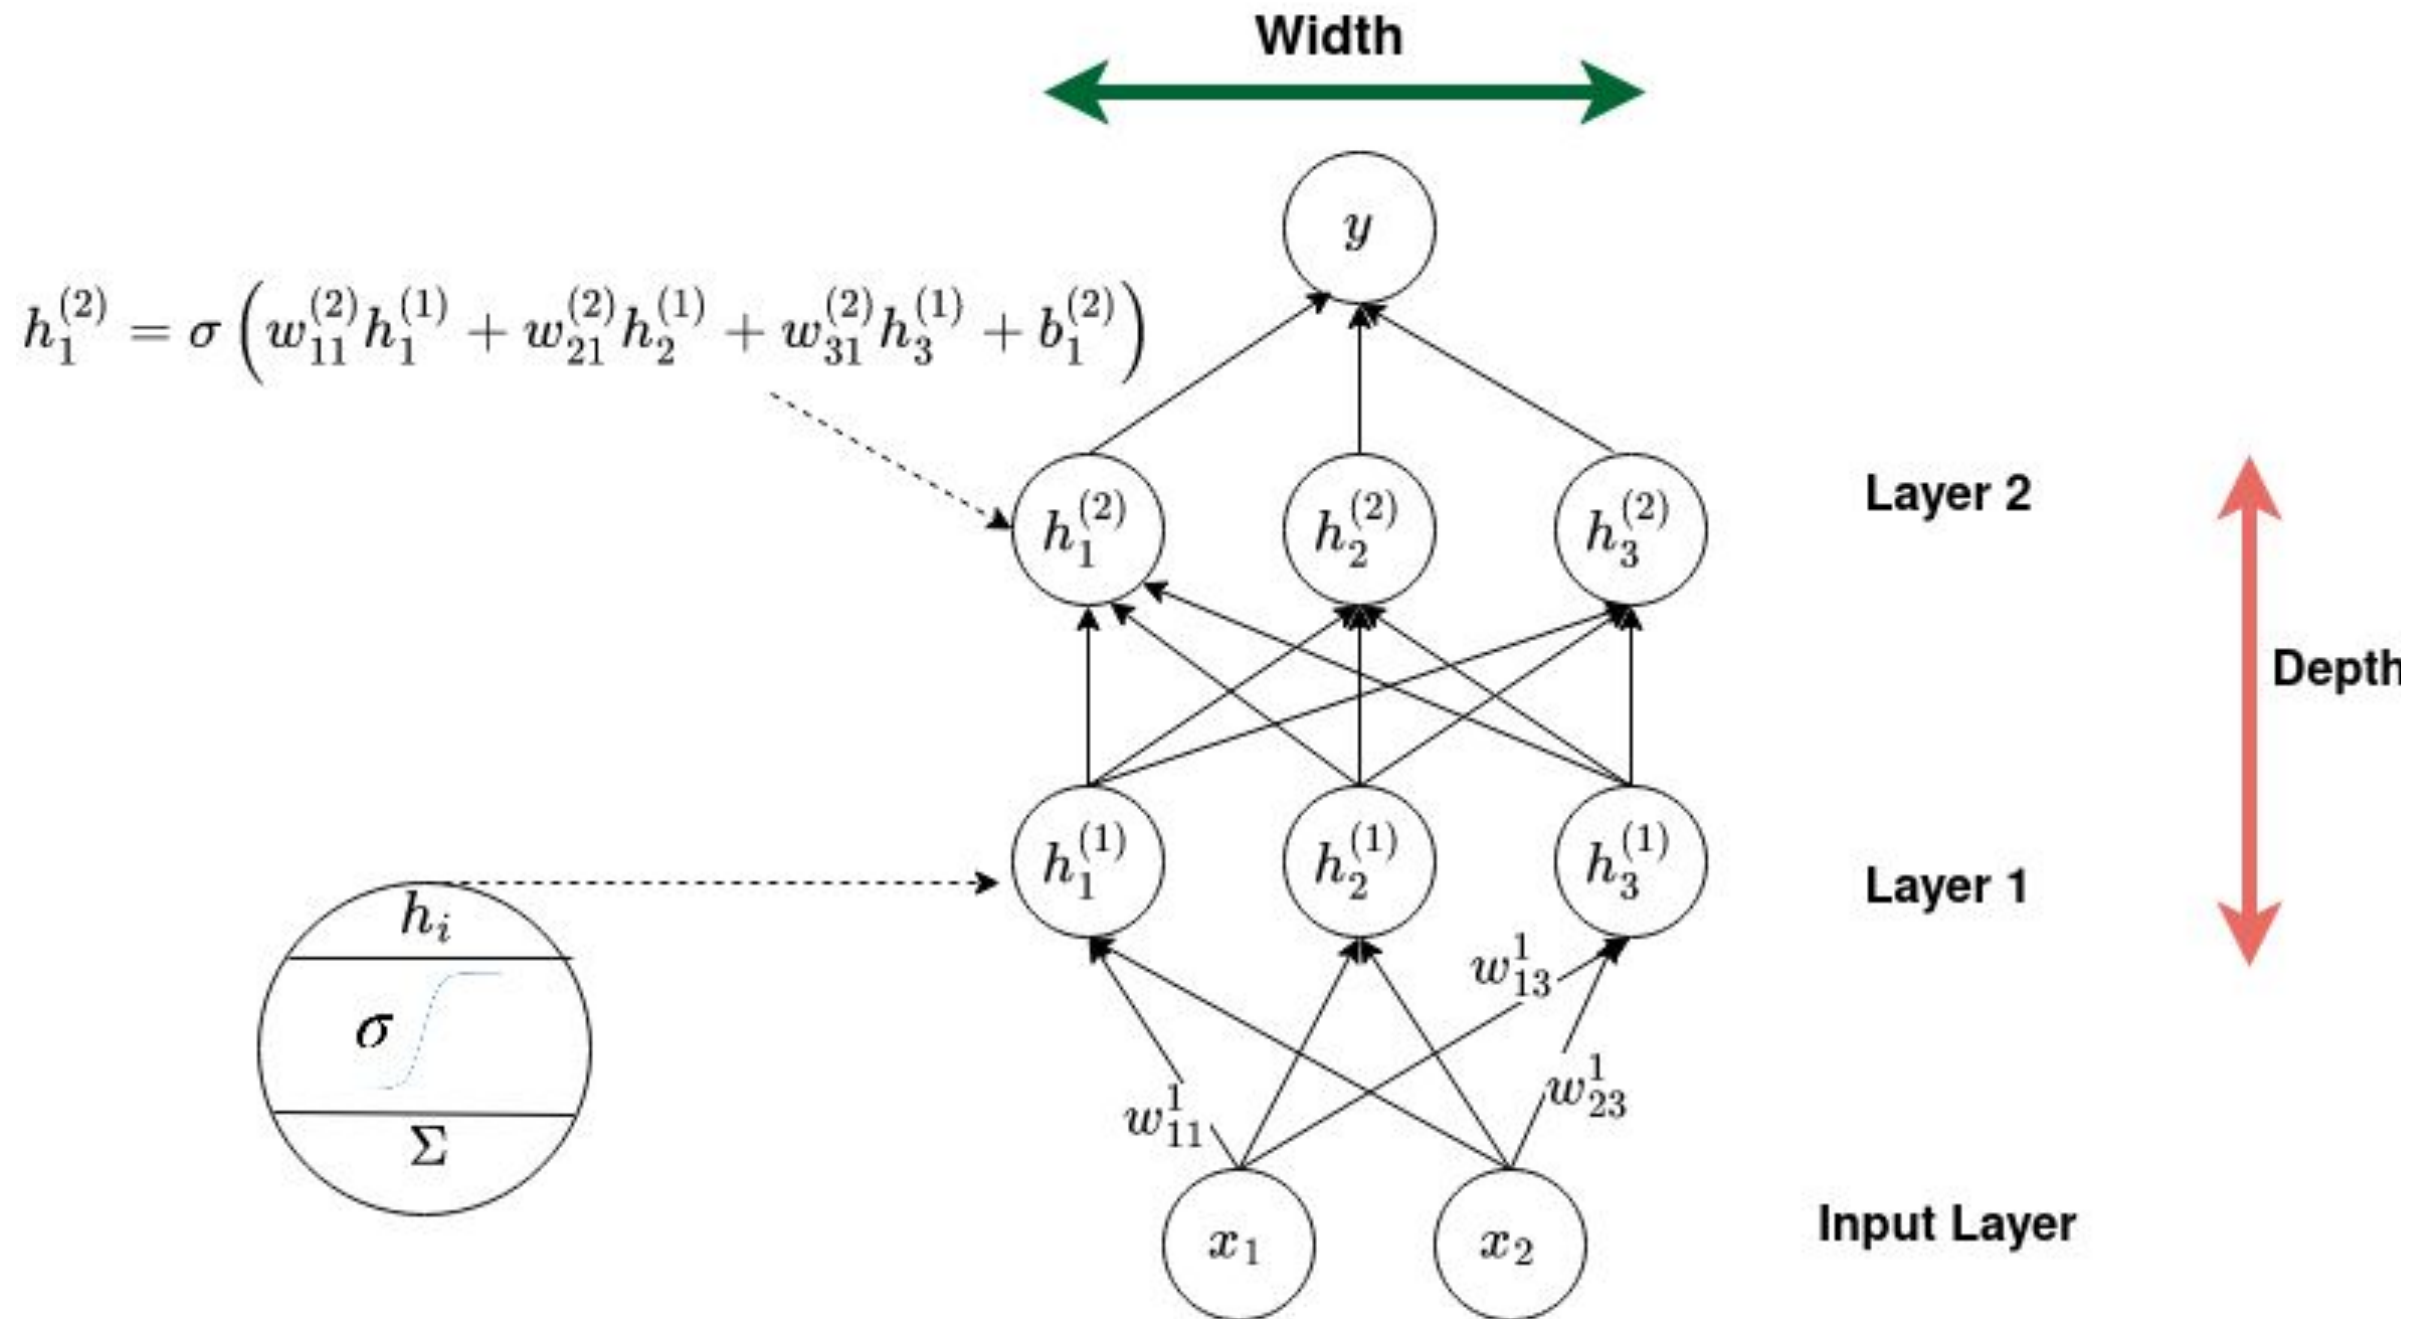
Neural Networks

- Simple computation blocks that work together



Neural Networks

- Simple computation blocks that work together



Review of Linear Algebra

- Scalar
- Vector
- Matrices
- Tensors

Representation in numerical software packages

Numpy

```
1 import numpy as np
2
3 # scalar
4 a = 5
5
6 # vector
7 b = np.array([1,1])
8 c = np.array([2,2])
9 np.dot(b,c)
```

4

```
1 A = np.array([[1,2],[3,4]])
2 B = np.array([[2,3],[4,1]])
3 print(A)
```

```
[[1 2]
 [3 4]]
```

PyTorch

```
1 import torch
2
3 b = torch.from_numpy(b)
4 c = torch.from_numpy(c)
5 print(b)
6
```

```
tensor([1, 1])
```

```
1 torch.dot(b,c)
```

```
tensor(4)
```


Matrix & Elementwise Multiplications

Numpy

```
1 import numpy as np
2 A = np.array([[1,2],[3,4]])
3 B = np.array([[1,1],[1,1]])
4
5 #Matrix multiply AB
6 print('Matrix Multiply:')
7 print(np.matmul(A,B))
8
9 print('ElementWise Multiply:')
10 #Elementwise Multiply
11 print(A*B)
```

```
Matrix Multiply:
[[3 3]
 [7 7]]
ElementWise Multiply:
[[1 2]
 [3 4]]
```

PyTorch

```
1 import torch
2 A = torch.tensor([[1,2],[3,4]])
3 B = torch.tensor([[1,1],[1,1]])
4
5 #Matrix multiply AB
6 print('Matrix Multiply:')
7 print(torch.matmul(A,B))
8
9 print('ElementWise Multiply:')
10 #Elementwise Multiply
11 print(A*B)
```

```
Matrix Multiply:
tensor([[3, 3],
        [7, 7]])
ElementWise Multiply:
tensor([[1, 2],
        [3, 4]])
```


Tensors

✓
0s

```
[36] 1 c = torch.tensor([[1,2],[3,4]])  
      2 print(c)  
      3 print(c.shape)
```

```
tensor([[1, 2],  
        [3, 4]])  
torch.Size([2, 2])
```

✓
0s

```
[37] 1 d = torch.tensor([[[1,2],[3,4]],  
                          2 [[3,2],[6,7]])]  
      3 print(d)  
      4 print(d.shape)
```

```
tensor([[[1, 2],  
         [3, 4]],  
        [[3, 2],  
         [6, 7]])]  
torch.Size([2, 2, 2])
```

Indexing

```
✓ [16] 1 d = torch.tensor([[[1,2],[3,4]],  
0s      2                [[3,2],[6,7]])  
      3 print(d)  
      4 print(d.shape)  
  
      tensor([[[1, 2],  
               [3, 4]],  
              [[3, 2],  
               [6, 7]])  
      torch.Size([2, 2, 2])
```

```
✓ [17] 1 print(d[0,:,:])  
0s  
  
      tensor([[[1, 2],  
               [3, 4]])
```

```
✓ [18] 1 print(d[0,0,:])  
0s  
      tensor([1, 2])
```

Matrix Inverse

Identity Matrix

$$\mathbf{I}_n \longrightarrow \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

Inverse

$$\mathbf{A}^{-1} \mathbf{A} = \mathbf{I}_n.$$

Norms

L_p norm

$$||\boldsymbol{x}||_p = \left(\sum_i |x_i|^p \right)^{\frac{1}{p}}$$

Shorthand for p=2, length of a vector


$$||\boldsymbol{x}||$$

Eigenvalues/vectors

- Eigenvalues are a property of a square matrix


Eigenvector is a vector such that applying A is a rescaling

$$Av = \lambda v.$$




Neural Networks

- Simple computation blocks that work together




Neural Networks

- Can be written as composition of simple linear algebra operations (matrix multiplies and elementwise non-linearities)



Neural Networks

- Can be written as composition of simple linear algebra operations (matrix multiplies and elementwise non-linearities)
- Bias term can be included in W to simplify notations





W_i Matrix of parameters at layer i
 σ Pointwise non-linearity

Intro to Machine Learning

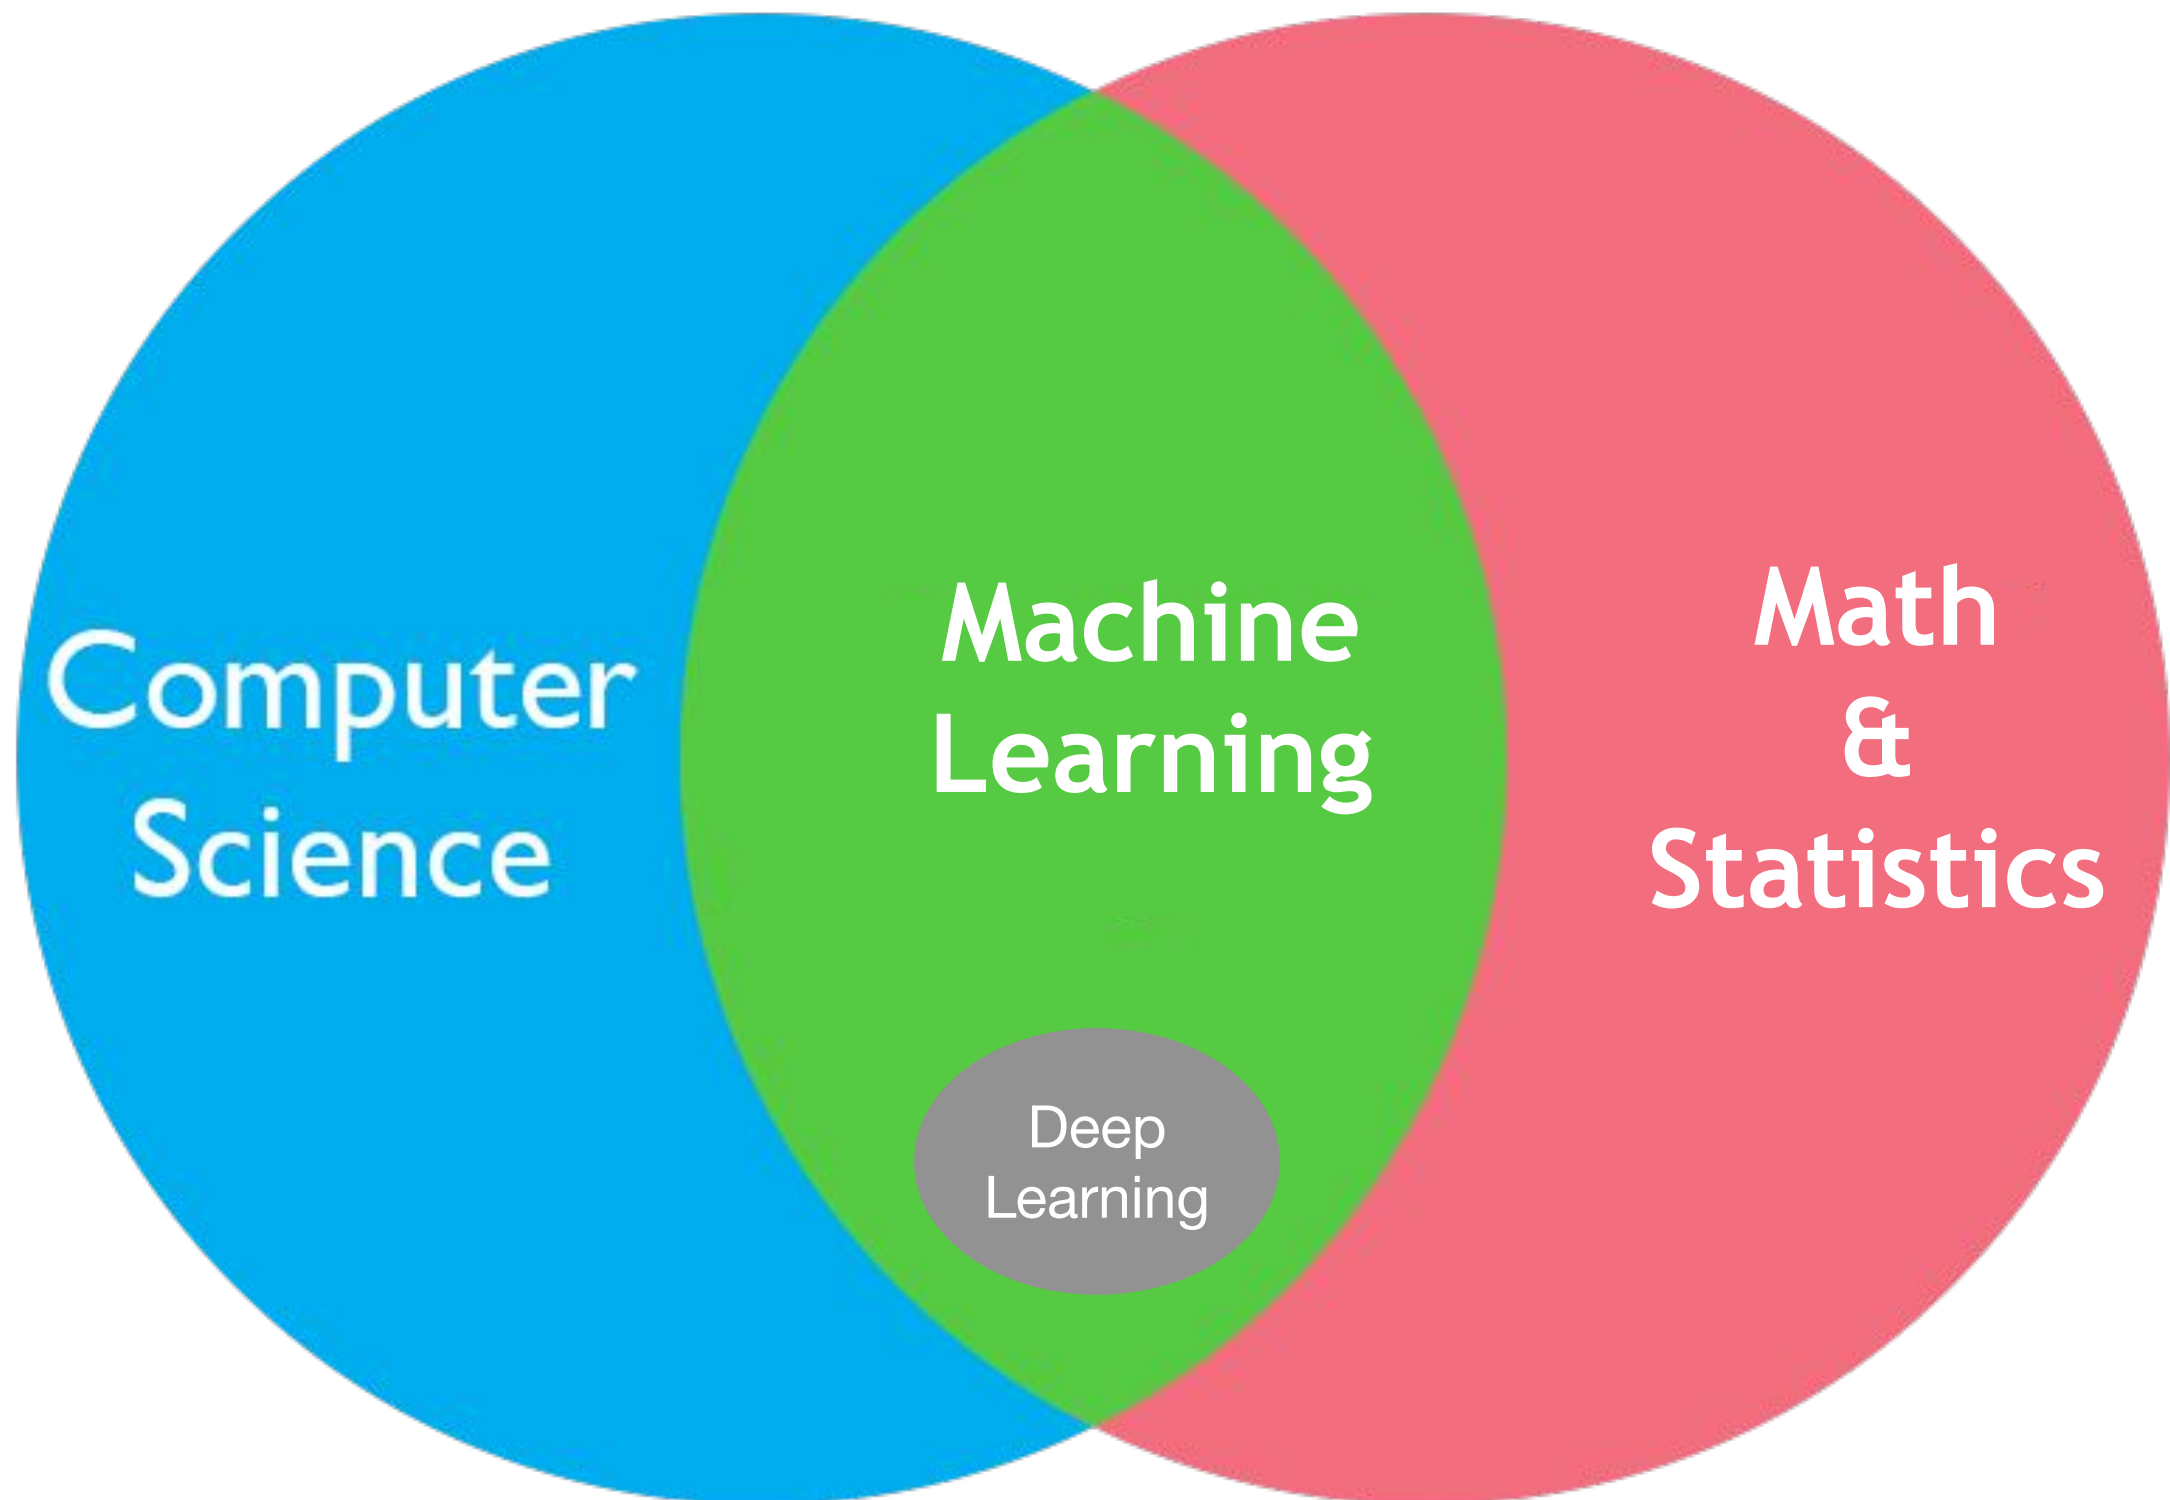
Data Explosion



- Millions and billions of measurements
- Millions and Billions of examples




What is Machine Learning




What is Machine Learning

- “Machine Learning: Field of study that gives computers the ability to learn without being explicitly programmed.” -Arthur Samuel (1959)
- “A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P , if its performance at tasks in T , as measured by P , improves with experience E .” Michel (1997)

Traditional Programming



Machine Learning



Types of Machine Learning

- **Supervised Learning** (Given: Training Data and Output examples)




Or



- **Unsupervised Learning** (Given: Training Data and no output)

- Data exploration and visualization



- Generating data



- **Reinforcement Learning**

- Rewards from sequences of actions

agent




Supervised Learning: Regression



How tall is my dog? (depending on age)

	Age	Height
Dog 1	1	3
Dog 2	3	4.5
Dog 3	1	1
Dog 4	5	5.5
Dog 5	9	16
Dog 6	4	5.2
Dog 7	6	7.5
Dog 8	12	20

Height



Supervised Learning: Regression

Training Data

	Age	Height
Dog 1	1	3
Dog 2	3	4.5
Dog 3	1	1
Dog 4	5	5.5
Dog 5	9	16
Dog 6	4	5.2
Dog 7	6	7.5
Dog 8	12	20

Use Model:

$$y = a * x + b$$


Optimize Objective:

$$\min_{a,b} \sum_i (\text{height}_i - (a * \text{age}_i + b))^2$$

Testing Data


	Age	Height
New Dog 1	1.2	?
New Dog 2	3.1	?

Height





Age


Overfitting and Underfitting



Underfit



Overfit




Supervised Learning: Classification

Classify cats and dog using height and weight



	Height	Weight
Cat 1	110 cm	20 kg
Cat 2	108 cm	15kg
Cat 3	109 cm	21kg
Cat 4	112.5 cm	23 kg
Dog 1	130 cm	28kg
Dog 2	125 cm	29kg
Dog 3	109.5 cm	30 kg
Dog 4	130 cm	22kg



Learning in High Dimensions Data is Harder

Classify cats and dog from raw pixels




VS




High Dimensional

Images are $3 \times 512 \times 512 \sim 1$ million dimensions



Non-Linear Boundary

Autonomous Driving Example



- Supervised Learning
 - Identify objects
 - Uncertainty
 - Driving decisions from example



- Unsupervised Learning
 - Which data to label?
- Reinforcement Learning
 - Improve driving decision
- Simulations